

NETWORK ROUTING PROTOCOL FAILURE CASE STUDY A

Root-Cause Analysis of the March 2024 NimbusEdge Network Outage

SUBJECT AREA

STUDENT	Rithin Sai
REGISTER NUMBER	1972260016
	Computer Networks

Based on the Incident Report, Network Topology Diagram, RCA Report, NOC Change Log and Client SLA Breach Notice.

CONTENTS

RIP — Chennai HQ Campus	Overall Conclusion
-------------------------	--------------------

4

IGRP — Bengaluru–Pune–Hyderabad
Region

8

5

Documents Referenced

OSPF — Core Backbone Migration

6

BGP — ISP Peering and Failover

8

7

CASE SNAPSHOT

The March 2024 outage affected six enterprise clients for approximately nine hours. The investigation identified four separate routing-protocol issues: a RIP hop-count limitation, slow IGRP convergence, an OSPF Area ID mismatch, and a BGP failover-policy problem.

01 | INTRODUCTION

In March 2024, NimbusEdge Technologies experienced a major network outage that disrupted service for six enterprise clients for approximately nine hours. Investigation showed that the outage was not the result of a single fault, but of four separate routing-protocol issues occurring across different parts of the network: a hop-count limitation in RIP, a slow convergence response in IGRP, a configuration mismatch in OSPF, and a failed automatic failover in BGP.

This report analyses each failure individually — what happened, why it happened, why it was difficult to detect, and what corrective action is recommended — before drawing an overall conclusion about the network’s design and operational practices.

02 | INCIDENT OVERVIEW

			marked unreachable Degraded link was not
	RIP	Bengaluru – Pune – Hyderabad Core	abandoned quickly enough (slow convergence)
IGRP		Backbone – Hyderabad ISP-A / ISP-B	Incorrect Area ID prevented adjacency
OSPF		Peering	from forming Automatic failover to the
BGP			backup ISP did not trigger as designed
		Delivery-department route reached 16 hops and was	
Chennai HQ			

03 | RIP — CHENNAI HQ CAMPUS

What happened?

Chennai HQ added a new Delivery-department subnet while the site was still running RIP v2. The route to this subnet reached a hop count of 16. Since RIP only treats routes of 1–15 hops as reachable, the Delivery subnet was silently marked unreachable. The fault surfaced only after client support tickets were raised because standard monitoring did not flag the routing failure.

Why did 16 hops cause the failure?

RIP uses hopcount — the number of routers a packet must cross — as its sole routing metric. Values of 1–15 are considered reachable; a metric of 16 represents infinity, meaning unreachable. Once the Delivery route's hop count reached 16, RIP stopped advertising it as usable.

Why was it not automatically detected?

This was a logical routing failure rather than a physical link failure: the routers and interfaces continued to function normally, so link-level monitoring showed no problem. The RCA identified a monitoring gap — route reachability was not tracked with the same rigour as physical link status — compounded by RIP's inherent 15-hop ceiling.

RECOMMENDED FIX

Migrate Chennai HQ from RIP to OSPF. OSPF scales further and has no hop-count ceiling, making it a better fit

as the site grows. RIP was a reasonable original choice for a small, stable campus because of its simplicity; the problem emerged once the network expanded beyond that scope.

RULE OF THUMB: 1–15 hops = reachable | 16 hops = unreachable

04 | IGRP — BENGALURU–PUNE–HYDERABAD

What happened?

An ISP-side fault degraded the Bengaluru–Pune fibre link: effective bandwidth dropped to roughly 2 Mbps and latency increased. A healthier alternate path existed via Pune–Hyderabad–Bengaluru, but IGRP took too long to shift traffic onto it, causing severe application latency for clients on the affected segment.

What does the IGRP metric show?

Classic IGRP computes a composite metric weighted mainly by bandwidth and delay. For a 2,000 kbps link, the bandwidth component works out to $10,000,000 \div 2,000 = 5,000$ before delay is added and final scaling applied. Because delay also rose, the composite metric for the degraded path should have worsened further, and the healthier alternate should have been preferred once its complete metric was lower.

Data limitation: The case materials do not provide complete delay values for both paths, so an exact side-by-side metric comparison would require assuming figures not present in the source documents.

Why was the problem mainly convergence?

IGRP's metric design was not itself at fault —using bandwidth and delay together is more informative than hop count alone. The real problem was speed of reaction: the network did not converge onto the healthier path quickly enough after conditions changed. In short, IGRP had the right information but acted on it too slowly.

RECOMMENDED FIX

Migrate the regional network to OSPF. IGRP was originally an improvement over RIP because it factored in bandwidth and delay, but a growing regional network needs the faster convergence and stronger scalability that OSPF provides, alongside active monitoring of both link performance and routing convergence time.

05 | OSPF — CORE BACKBONE MIGRATION

What happened?

NimbusEdge was migrating its core network to multi-area OSPF, with Hyderabad intended to sit in Area 3. The Hyderabad router was instead configured with the wrong Area ID. Because the ID did not match the design, the router could not form the expected OSPF adjacency, creating a routing black hole until a later change record corrected the Area ID.

Why does the Area ID matter?

OSPF neighbours must agree on several parameters before forming an adjacency: Area ID, Hello interval, Dead interval, authentication settings (where used), and compatible network types. In this incident the specific mismatch was the Area ID — because the two routers disagreed on which area they belonged to, the neighbour relationship could not form at all.

Why was it harder to detect?

A physical link failure is usually obvious, since the interface itself goes down. Here the physical connection stayed up while the routing layer silently failed, so basic link monitoring showed everything as healthy. The RCA also points to a change-management gap: the misconfiguration was not caught because post-change validation was insufficient.

HOW SHOULD IT BE FIXED?

- Correct the Hyderabad router's Area ID to the intended Area 3 assignment.
- Verify that the OSPF neighbour reaches FULL state.

- Confirm route/LSA exchange and check the routing table.
- Run an end-to-end connectivity test.
- Introduce peer review, a formal change checklist, and mandatory post-change adjacency verification.

OSPF offers faster convergence, better scalability, and detailed link-state visibility compared with the older RIP/IGRP design, making it the more appropriate protocol for a multi-site enterprise backbone.

06 | BGP — ISP PEERING AND FAILOVER

What happened?

NimbusEdge maintained two ISP connections: ISP-A as primary and ISP-B as backup. When ISP-A failed, BGP was expected to withdraw routes through ISP-A and automatically select ISP-B. This automatic failover did not occur, forcing engineers to intervene manually — a delay that contributed directly to the length of the outage and to the resulting SLA breach.

Key BGP attributes

LOCAL_PREF	AS-PATH	— higher value is preferred. The sequence of autonomous systems a
		route has traversed; a shorter path is generally preferred wh Influences the
MED		preferred entry point between two connected autonomous systems; a
	Preference for an outbound path within the company's own AS	lower value is ge

The RCA points specifically to a LOCAL_PREF / AS-PATH policy problem as the reason automatic failover did not behave as expected.

Why is an untested BGP policy risky?

BGP governs communication with externalautonomoussystems and ISPs, so a policy mistake can affect a large volume of external traffic. The resulting chain was: ISP-A failed → automatic failover did not trigger → engineers intervened manually → recovery was delayed → clients were impacted → the SLA was breached.

What should the BGP policy look like?

During normal operation, ISP-A should carry a higher preference than ISP-B — for example, LOCAL_PREF 200 for ISP-A versus 100 for ISP-B. These are illustrative values; the key requirement is that ISP-A remains primary and ISP-B remains backup. When ISP-A's route is withdrawn, ISP-B should automatically become the best available path. NimbusEdge should also run regular, controlled failover tests.

automatic switch to ISP-B.

ISP-A failed, but the existing policy did not produce the expected

ISP-A is primary with higher LOCAL_PREF; ISP-B is backup with lower LOCAL_PREF. When ISP-A fails, its route is withdrawn and ISP-B becomes active automatically.

07 | OVERALL CONCLUSION

The NimbusEdge outage was the product of four distinct routing failures rather than one isolated cause: RIP hit its inherent 15-hop limit, IGRP reacted too slowly to a degraded link, OSPF was undermined by an incorrect Area ID, and BGP failed to deliver the automatic ISP failover it was designed for.

The common thread

Effective routing depends on more than protocol selection. It requires correct configuration, meaningful monitoring, disciplined change validation, and regular, realistic failure testing.

	for the network's	correctly.	TEST
	scale.	MONITOR	failover before a
CHOOSE	CONFIGURE	routing state, not	real incident.
the right protocol	routing parameters	just link status.	

In simple terms: choose the right protocol for the network's scale, configure it correctly, monitor the actual routing state (not just link status), and test what happens when a link or path fails — before a real incident forces the test.

08 | DOCUMENTS REFERENCED

- Incident_Report_NimbusEdge_March2024.pdf

- Network_Topology_Diagram_NimbusEdge.pdf
- RCA_Report_Q1_2024.docx
- NOC_Change_Log_Feb–Mar_2024.xlsx
- Client_SLA_Breach_Notice.pdf